

PromptGFM-Bio: Project Progress Report

Gene Prioritization for Rare Diseases via Prompt-Conditioned Graph Foundation Model

Prepared for: Faculty Advisor

Date: April 2026

Status: Ablation experiments in progress — publication targeting NeurIPS/ICML 2026 Bio-ML Workshop

1. Problem Statement

Rare diseases affect approximately 300 million people worldwide, yet the majority have no effective treatment. A fundamental bottleneck is gene discovery: identifying which genes are causally linked to a disease phenotype. For most rare diseases, only 1–5 causal genes are known, making supervised learning approaches difficult and motivating the need for models that can generalize to diseases with little or no prior genetic knowledge.

The core task this project addresses is gene prioritization: given a disease and its clinical description, rank all ~19,576 candidate genes by their likelihood of being causally associated with that disease. This is formulated as a link prediction problem on a heterogeneous biological knowledge graph.

The central research question is:

Can dynamically conditioning graph neural network message passing on natural language disease descriptions — via Feature-wise Linear Modulation (FiLM) — improve gene prioritization for rare diseases, including zero-shot generalization to diseases never seen during training?

2. Why This Problem is Hard

Standard supervised approaches fail here for three reasons:

Extreme class imbalance. Each disease has 2–5 known gene associations out of ~19,576 candidates — a positive rate of ~0.01%. This makes precision-oriented metrics inherently low and requires careful evaluation design.

Sparse labels for rare diseases. Ultra-rare diseases (≤ 5 known gene associations) cannot be learned from labeled data alone. A model must generalize from disease semantics.


Static feature concatenation is insufficient. Prior work typically encodes disease text into a feature vector and concatenates it with graph embeddings. This gives every gene the same disease-agnostic representation during message passing — the language signal only enters at the prediction head, not during graph reasoning itself.

3. Proposed Solution: PromptGFM-Bio


PromptGFM-Bio introduces prompt-conditioned message passing: the disease's natural language description dynamically modulates the GNN's feature transformation at every layer via FiLM conditioning. This means different disease prompts produce different gene representations from the same graph — the model learns to "look at the PPI network differently" for each disease.


3.1 Architecture Overview

Disease Text Description



↓ prompt embedding $p \in \mathbb{R}^{768}$





3.2 Component Details

Prompt Encoder

- Model: microsoft/BiomedNLP-PubMedBERT-base-uncased-abstract-fulltext
- Pooling: CLS token → 768-dimensional disease embedding
- Frozen during training (no fine-tuning of BERT weights)
- Embeddings pre-computed and cached for all 12,714 diseases before training begins — BERT is never called per-batch, making training fast

GNN Backbone: GraphSAGE

- 3 message-passing layers
- Dimensions: 128 → 512 → 512
- Operates on heterogeneous graph with gene-gene PPI edges
- Aggregation: mean pooling of neighborhood features

FiLM Conditioning

- For each GNN layer output h : $\gamma, \beta = \text{MLP}(p)$; $h' = \gamma \odot h + \beta$
- γ (scale) and β (shift) are learned functions of the disease prompt
- Initialized so $\gamma \approx 1$, $\beta \approx 0$ (identity at start of training)
- Key property: same PPI graph produces different gene embeddings for different diseases

Predictor

- Input: concatenation of conditioned gene embedding + disease embedding
- Architecture: MLP with hidden_dim=256
- Output: scalar score per (gene, disease) pair

3.3 Training Objective

Combined loss function with three components:

$$L = 1.0 \times L_{\text{BCE}} + 0.5 \times L_{\text{ranking}} + 0.3 \times L_{\text{listnet}}$$

- L_{BCE} : Binary cross-entropy on positive/negative pairs
- L_{ranking} : Pairwise ranking loss (margin=0.5) ensuring positives rank above negatives
- L_{listnet} : ListNet loss for list-level ranking optimization
- Negative sampling: 5 random negatives per positive edge during training

Optimizer: AdamW (lr=5e-4, weight_decay=0.01, cosine schedule, 5 warmup epochs) Mixed precision: AMP (FP16) enabled — 1.5–2× speedup Batch size: 768 (auto-scaled to available VRAM)

4. Knowledge Graph

4.1 Data Sources

Source	Type	Content	File
STRING v12	PPI network	Protein-protein interactions, human (9606)	9606.protein.links.v12.0.txt.gz
BioGRID 4.4.224	PPI network	Experimentally validated interactions	BIOGRID-ALL-4.4.224.tab3.zip
DisGeNET	Gene-disease	Curated gene-disease associations	curated_gene_disease_associations.tsv
HPO	Phenotype	Gene-phenotype mappings, disease ontology	phenotype_to_genes.txt, phenotype.hpoa
Orphanet	Rare disease	Rare disease gene associations	en_product1.xml, en_product4.xml, en_product6.xml

4.2 Graph Statistics

Entity	Count
Gene nodes	19,576
Disease nodes	16,841
Phenotype nodes	11,794
Gene ↔ Gene edges (STRING PPI ≥700)	1,854,012
Gene → Disease edges	9,741,610
Disease → Gene edges (reverse)	9,741,610

PPI threshold: STRING confidence score ≥700 out of 1000 (high-confidence interactions only)

4.3 Edge File for Supervised Learning

The training labels come from hpo_gene_disease_edges.csv (HPO phenotype bridge):

- Raw: 9,734,247 edges
- After score filter (≥0.3): 1,170,143 edges
- Unique genes in vocabulary: 5,251
- Unique diseases in vocabulary: 12,714

4.4 Train/Val/Test Split

Split	Edges	Notes
Train	936,114 (80%)	Used for gradient updates
Validation	117,014 (10%)	Early stopping on Val AUROC
Test	117,015 (10%)	Final evaluation only

Split is edge-level random with seed=42. The same split is reproduced exactly for all ablation runs.

5. Evaluation Protocol

5.1 Full-Vocabulary Ranking

All evaluations rank every gene in the vocabulary (19,576 genes) for each disease query. This is a deliberate design choice that produces honest, publication-ready metrics. For each disease in the test set:

1. Score all 19,576 candidate genes using the model
2. Sort by descending score
3. Compute ranking metrics against known true associations

This contrasts with the naive approach of ranking 1 positive against 5 random negatives, which would give artificially inflated Precision@K values.

5.2 Metrics

Metric	Formula	What it measures
AUROC	Area under ROC curve	Global discrimination ability
AUPR	Area under Precision-Recall	Discrimination at high precision (note: inherently low at 19,576-gene ranking)
Hit Rate@K	% of queries where true gene is in top K	Clinical utility — is the right answer near the top?
MRR	Mean(1/rank of first true gene)	Average reciprocal rank of true associations
NDCG@K	Normalized discounted cumulative gain	Ranking quality with position weighting
Precision@K	True positives in top K / K	Precision at cutoff K
Recall@K	True positives in top K / all positives	Recall at cutoff K
MAP	Mean average precision	Average precision across all queries

K values evaluated: 10, 20, 50, 100

Primary paper metrics (in order of importance):

4. Hit Rate@50 — most clinically meaningful
5. AUROC — best headline number
6. MRR — summary of ranking quality

6. Results

6.1 Full Model Performance (v6 — Current Best)

Trained 100 epochs on RTX 4090, best checkpoint at epoch 87.

Metric	Value
Test AUROC	0.9626
Test AUPR	0.0188
Hit Rate@10	0.3218
Hit Rate@20	0.4175
Hit Rate@50	0.5573
Hit Rate@100	0.6537
MRR	0.1588

Metric	Value
MAP	0.0569
NDCG@10	0.0795
NDCG@50	0.1137
NDCG@100	0.1384
Recall@50	0.1941
Recall@100	0.2801

Interpretation: For 55.7% of rare disease queries, the true causal gene appears in the top 50 out of 19,576 candidates — a 392× enrichment over random chance. For 1 in 3 diseases, the true gene is in the top 10 (322× enrichment).

The low AUPR (0.0188) is mathematically expected and not a flaw: with ~2 true positives out of 19,576 candidates (0.01% positive rate), even a perfect ranker would have low AUPR. This is the correct, honest evaluation.

6.2 Training History (v6)

Phase	AUROC	Key event
Epoch 1	0.9243	Baseline after 1 epoch
Epoch 20	0.9400	Rapid early learning
Epoch 50	0.9510	Continued improvement
Epoch 87	0.9547	Best checkpoint saved
Epoch 100	0.9543	Converged, no improvement for 13 epochs

7. Ablation Study (In Progress)

7.1 Design

Four variants isolate the contribution of each architectural component:

Variant	use_gnn	use_conditioning	Description
MLP only	✗	✗	Raw node embeddings → MLP predictor. No message passing, no prompt conditioning. Pure baseline.
Prompt only	✗	✓	FiLM conditioning active, but no PPI edges → GNN has no neighbors to aggregate. Prompt signal enters at prediction only.
GNN only	✓	✗	PPI message passing active, FiLM identity ($\gamma=1$, $\beta=0$). GraphSAGE without language guidance.
Full PromptGFM-Bio	✓	✓	Complete model. PPI message passing dynamically conditioned on disease text.

Each variant is trained with 3 random seeds (42, 43, 44) and results reported as mean $\pm$ std.

7.2 Results So Far

Variant	Val AUROC (best)	Status
MLP only (seed 42)	0.9557 at epoch 97	✅ Complete
MLP only (seeds 43, 44)	TBD	✅ Complete
Prompt only (seed 42)	0.9435 at epoch 10+	🔄 Running
GNN only (×3 seeds)	TBD	⌚ Queued
Full model (×3 seeds)	0.9547 (from v6)	✅ Have result

Notable early finding: The MLP-only baseline converges at Val AUROC ~0.955, significantly higher than the pre-fix baseline (0.813) reported in early experiments. This is because the earlier evaluation used a small negative pool (5 negatives per positive) rather than ranking all genes — the evaluation was corrected as part of this project.

7.3 Expected Paper Table

Variant	AUROC	Hit@50	MRR
MLP only	~0.955	TBD	TBD
Prompt only (FiLM, no PPI)	~0.956	TBD	TBD
GNN only (PPI, no FiLM)	TBD	TBD	TBD
Full PromptGFM-Bio	0.9626	0.5573	0.159

8. Zero-Shot Rare Disease Evaluation

8.1 Motivation

A key claim of the paper is zero-shot generalization: given only a disease's text description, can the model prioritize genes for diseases it has never seen any gene associations for during training?

8.2 Zero-Shot Disease Identification

Using scripts/find_rare_diseases.py, we identify diseases with ≤ 5 known gene associations that appear in neither training nor validation splits:

Category	Count	% of dataset
Total diseases	12,714	100%
≤ 5 gene associations (candidate)	1,871	14.7%
Excluded (appear in training)	1,754	—
Clean zero-shot set	117	0.9%

The 117 zero-shot diseases span OMIM and Orphanet entries covering a range of rare genetic conditions including Angelman syndrome-related, X-linked disorders, and ultra-rare metabolic diseases.

This exceeds SHEPHERD (NeurIPS 2023), which used ~50 rare diseases for zero-shot evaluation.

8.3 Zero-Shot Protocol

For each of the 117 zero-shot diseases:

7. The model receives only the disease name/description as input (no training labels)

- 8. All 19,576 genes are ranked
- 9. Metrics: Hit@10, Hit@50, MRR

Zero-shot evaluation will be run after ablation training completes using:

```
python3 scripts/evaluate.py \
--config configs/ablations/ablation_4_full_model.yaml \
--checkpoint checkpoints/ablation_4_full_model/best_model.pt \
--zero_shot
```

9. Comparison to Related Work

9.1 What Differentiates PromptGFM-Bio

Approach	Method	Limitation
Standard GNN (Hamilton et al.)	Message passing only	No language signal
Feature concat (most prior work)	Concat text embedding to node features	Static — same representation regardless of disease
SHEPHERD (Bejan et al., 2023)	GNN + HPO phenotype graph	No dynamic language conditioning of message passing
BioKGBERT-style	BERT on KG triples	No graph message passing
PromptGFM-Bio	FiLM-conditioned GraphSAGE + frozen PubMedBERT	Dynamic: disease text modulates every message-passing step

9.2 Key Novelty Claim

The core contribution is not simply combining a GNN with a language model — that has been done. The contribution is where the language signal enters: not at the prediction head after the graph computation, but inside the message passing itself, via FiLM modulation. This enables the model to learn disease-specific graph traversal patterns.

10. Infrastructure and Reproducibility

10.1 Hardware

Component	Specification
Machine	ARC Labs Workstation 6
CPU	Intel i9-14900K
RAM	128 GB
GPU	NVIDIA GeForce RTX 4090 (24 GB VRAM)
CUDA	12.4 / Driver 580.65.06
Storage	512 GB

10.2 Software Stack

Component	Version
Python	3.12
PyTorch	2.6.0+cu124

Component	Version
PyTorch Geometric	Latest (GraphSAGE, HeteroData)
HuggingFace Transformers	≥4.40
Environment	micromamba promptgfm

10.3 Repository Structure

PromptGFM-Bio/

```

├── src/
│   ├── models/
│   │   ├── promptgfm.py    ← Main model (use_gnn + use_conditioning flags)
│   │   ├── gnn_backbone.py ← GraphSAGE encoder
│   │   ├── conditioning.py ← FiLM conditioning
│   │   └── prompt_encoder.py ← Frozen PubMedBERT
│   ├── training/
│   │   ├── finetune.py    ← Training loop (AMP, early stopping)
│   │   └── losses.py      ← BCE + ranking + listnet
│   ├── data/
│   │   ├── dataset.py     ← GeneDiseaseDataset, graph loading
│   │   └── preprocess.py  ← Graph construction from raw data
│   ├── evaluation/
│   │   └── metrics.py     ← Full-vocabulary gene ranking evaluator
│   ├── scripts/
│   │   ├── train.py       ← Entry point (fork-context patched)
│   │   ├── evaluate.py    ← Evaluation (+ --zero_shot flag)
│   │   ├── find_rare_diseases.py ← Identifies 117 zero-shot diseases
│   │   └── run_ablations.sh ← 4 configs × 3 seeds = 12 runs
│   ├── configs/
│   │   ├── workstation_config.yaml ← Active config
│   │   └── ablations/          ← 4 ablation YAML files
│   ├── checkpoints/
│   │   ├── promptgfm_film/    ← v6 best model (epoch 87)
│   │   └── ablation_*/       ← Per-ablation checkpoints
│   └── data/
│       ├── raw/              ← STRING, BioGRID, DisGeNET, HPO, Orphanet
│       ├── processed/
│       │   └── biomedical_graph.pt ← Built heterogeneous graph
│       └── splits/
│           └── zero_shot_rare_diseases.json ← 117 disease IDs

```

10.4 Training Speed

- ~93.5 seconds per epoch on RTX 4090
- 100 epochs ≈ 2.6 hours per run
- 12 ablation runs ≈ 31 hours total (running now, completes ~April 6 evening)

11. Known Issues Resolved During Development

Issue	Root Cause	Fix
Broken evaluation metrics	Precision@10=1.0, Recall@10=0.0001 simultaneously — ranking against only 5 negatives	Rewrote evaluator to rank all 19,576 genes per query

Issue	Root Cause	Fix
No message passing	Graph had only gene ↔ disease edges; no gene-gene PPI edges	Added 1,854,012 STRING PPI edges (≥700 confidence)
DataLoader Bus Error	Graph tensors sent through every batch → /dev/shm exhaustion	Store graph tensors once on trainer; use multiprocessing_context='fork'
KeyError: 'node_features'	Batch contract changed without updating trainer + validator	Synced all three components together
AMP deprecation warning	torch.cuda.amp.autocast deprecated in PyTorch 2.6	Updated to torch.amp.autocast('cuda', ...)

12. Publication Plan

12.1 Target Venues

Primary (NeurIPS 2026, deadlines August–September 2026):

Workshop	Fit	Reason
Learning on Graphs (LoG)	★★★★	Direct fit: GNN architecture + novel conditioning
Graph Learning for Drug Discovery	★★★★	Most relevant: gene-disease, biological graphs
AI for Science	★★★★	Broad bio-ML, well-attended

Backup (ICML 2026, deadlines April–May 2026):

- Computational Biology workshop
- Machine Learning for Genomics

12.2 Paper Format

Standard 4-page workshop paper + references. Sections:

10. Introduction + motivation
11. Related work
12. Methods (architecture, training)
13. Experiments (ablation table + zero-shot table)
14. Case studies (3 rare diseases, top-5 predicted genes)
15. Conclusion

12.3 Remaining Work

Task	Status	Timeline
Ablation runs (12 total)	🔄 Running	Complete April 6
Evaluate all 4 ablation models	🕒	April 7
Zero-shot evaluation	🕒	April 7
Literature review + related work	🕒	This week
Methods section draft	🕒	April 8–10
Results section + tables	🕒	April 10–12
Abstract + intro	🕒	April 12–14

Task	Status	Timeline
Adversarial reviewer review		April 14
Final polish + submission prep		May–September 2026

13. Summary

PromptGFM-Bio is a working end-to-end system for rare disease gene prioritization that:

16. Constructs a heterogeneous biological knowledge graph from STRING PPI, DisGeNET, HPO, and Orphanet — 19,576 genes, 16,841 diseases, 1.85M interaction edges
17. Trains a prompt-conditioned GraphSAGE model where disease text dynamically modulates message passing via FiLM — not just at the prediction head
18. Achieves AUROC=0.9626 and Hit Rate@50=55.7% on full-vocabulary gene ranking across 10,267 test disease queries
19. Identifies 117 clean zero-shot rare diseases for evaluating generalization to unseen conditions
20. Has a clear ablation study design (4 variants × 3 seeds = 12 runs) that isolates the contribution of PPI message passing vs. FiLM conditioning independently
21. Targets NeurIPS/ICML 2026 workshop venues with a differentiated novelty claim around dynamic, prompt-conditioned graph reasoning

The key finding emerging from early ablation results is that the MLP baseline is itself strong (~0.955 AUROC), suggesting the gene-disease association signal in the HPO/DisGeNET data is rich. The added value of PPI message passing and FiLM conditioning will be quantified precisely when the 12 ablation runs complete (estimated: April 6, 2026).

Repository: github.com/pes1ug23am910/PromptGFM-Bio (private) Compute: ARC Labs, AIML Department, Workstation 6 (RTX 4090)